

- import cPickle as pickle in order to transparently use the same interface.
- import job_stats from the tacc_stats monitor directory

```
In [1]: import sys
sys.path.append('../..//monitor')
import job_stats
import cPickle as pickle
```

- Load a single job's file from my directory of them on Ranger
- This is a local untar of John's nightly files

```
In [2]: j=pickle.load(open('nightly_jobs/2012-10-30/2887373'))
```

- Let's see what's actually in this object.

```
In [3]: dir(j)
```

```
Out[3]: ['__class__',
         '__delattr__',
         '__doc__',
         '__format__',
         '__getattribute__',
         '__hash__',
         '__init__',
         '__module__',
         '__new__',
         '__reduce__',
         '__reduce_ex__',
         '__repr__',
         '__setattr__',
         '__sizeof__',
         '__slots__',
         '__str__',
         '__subclasshook__',
         'acct',
         'aggregate_stats',
         'end_time',
         'error',
         'gather_stats',
         'get_schema',
         'get_stats',
         'hosts',
         'id',
         'munge_times',
         'process_dev_stats',
         'process_stats',
         'schemas',
         'start_time',
         'times',
         'trace']
```

- The hosts data structure contains most of the actual data
- But we'll come back to a few others

```
In [4]: j.hosts
```

```
Out[4]: {'i101-111.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204eed0>,  
'i112-306.ranger.tacc.utexas.edu': <job_stats.Host at 0x10b2e10d0>,  
'i113-106.ranger.tacc.utexas.edu': <job_stats.Host at 0x10b2e1110>,  
'i113-108.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ee10>,  
'i119-101.ranger.tacc.utexas.edu': <job_stats.Host at 0x10b2e1090>,  
'i133-403.ranger.tacc.utexas.edu': <job_stats.Host at 0x101dbe3d0>,  
'i137-102.ranger.tacc.utexas.edu': <job_stats.Host at 0x101dbb8d0>,  
'i145-106.ranger.tacc.utexas.edu': <job_stats.Host at 0x101e5bb10>,  
'i149-102.ranger.tacc.utexas.edu': <job_stats.Host at 0x101dbe410>,  
'i149-311.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ef50>,  
'i152-401.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204edd0>,  
'i154-410.ranger.tacc.utexas.edu': <job_stats.Host at 0x101dbe510>,  
'i158-310.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ee50>,  
'i162-207.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ef10>,  
'i168-212.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204efd0>,  
'i169-211.ranger.tacc.utexas.edu': <job_stats.Host at 0x101dbbe10>,  
'i176-109.ranger.tacc.utexas.edu': <job_stats.Host at 0x10b2e1050>,  
'i176-112.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ee90>,  
'i181-306.ranger.tacc.utexas.edu': <job_stats.Host at 0x10204ef90>}
```

- Hosts is a dictionary of further objects.
- Let's pick one and have a look around

```
In [5]: h='i101-111.ranger.tacc.utexas.edu'  
dir(j.hosts[h])
```

```
Out[5]: ['__class__',  
         '__delattr__',  
         '__dict__',  
         '__doc__',  
         '__format__',  
         '__getattr__',  
         '__hash__',  
         '__init__',  
         '__module__',  
         '__new__',  
         '__reduce__',  
         '__reduce_ex__',  
         '__repr__',  
         '__setattr__',  
         '__sizeof__',  
         '__str__',  
         '__subclasshook__',  
         '__weakref__',  
         'error',  
         'gather_stats',  
         'get_stats',  
         'get_stats_paths',  
         'job',  
         'marks',  
         'name',  
         'parse_stats',  
         'read_stats_file',  
         'read_stats_file_header',  
         'stats',  
         'trace']
```

- The stats dictionary contains all of the data for this host

```
In [6]: j.hosts[h].stats.keys()
```

```
Out[6]: ['ps',  
         'amd64_sock',  
         'lnet',  
         'mem',  
         'llite',  
         'vm',  
         'sysv_shm',  
         'amd64_core',  
         'numa',  
         'net',  
         'tmpfs',  
         'vfs',  
         'cpu',  
         'block',  
         'ib_sw']
```

- Each key here corresponds (more or less) to the lines from the raw data
- The value for each key is a dictionary containing key/value pairs that point to numpy arrays of the actual data for this type.
- In the case of amd64_core, there's a dictionary entry for each core (strings '0' through '15' on Ranger)

```
In [7]: j.hosts[h].stats['amd64_core']['0']
```

```
Out[7]: array([[ 0, 0, 0],
 [ 672322486793, 244443211072, 23836512957],
 [ 1582905469054, 570433271488, 57033734781],
 [ 2630260890974, 772689561088, 72806955413],
 [ 3544111028919, 1056275726016, 94484478062],
 [ 4469708202217, 1339141086272, 116993344068],
 [ 5403724140262, 1622699744064, 140416884014],
 [ 6339758624511, 1907506659072, 164308891339],
 [ 7278085528861, 2164789184896, 184305633969],
 [ 8223626657934, 2433906842368, 205218923263],
 [ 9155480787591, 2710674718336, 226568975083],
 [ 10099406145982, 2991129969536, 249877419073],
 [ 11031855201532, 3253103970304, 270527680112],
 [ 11978492245378, 3522650667392, 291467237077],
 [ 12931241269100, 3792735238144, 313293140007],
 [ 13877399003702, 4071925403776, 336508041477],
 [ 14817008040545, 4333454289984, 357992006775],
 [ 15755694983207, 4607378583616, 379178111515],
 [ 16689877183254, 4886062543232, 401003272674],
 [ 17641752260134, 5166123192192, 424697788352],
 [ 18599052783044, 5439558904256, 447578226683],
 [ 19531135916388, 5697525862272, 467208752172],
 [ 20462672918290, 5974731499136, 488585464240],
 [ 21410756621530, 6253488031488, 511800133446],
 [ 22355623050282, 6510835389632, 533044554937],
 [ 23297734837795, 6783576448128, 554174603954],
 [ 24245171225946, 7058298770880, 576288959216],
 [ 25180973207471, 7342083936256, 599796891201],
 [ 26113161144089, 7606714535040, 620612696985],
 [ 27048810221292, 7882293906176, 641902657767],
 [ 27999162256959, 8160595163840, 665094588944],
 [ 28922810328250, 8430981315584, 686540868604],
 [ 29859976319564, 8708150395264, 708395167192],
 [ 30804794073287, 8983069847808, 730599101169],
 [ 31745121433107, 9254894721536, 752778153174],
 [ 32672271371667, 9530151017664, 774810459287],
 [ 33605178473514, 9807339687488, 796211334444],
 [ 34560034791350, 10081698405568, 819180713197],
 [ 35487618120062, 10347569215424, 840090212933],
 [ 36417962286177, 10624208840320, 861444700425],
 [ 37351781939016, 10898945002688, 882681660275],
 [ 38301911411267, 11174700467136, 905701392186],
 [ 39245779520018, 11453439199232, 928885469720],
 [ 40196607931819, 11724025806656, 951572953447],
 [ 41154326981092, 11971335575552, 970512446915],
 [ 42106465603804, 12239052142848, 991327331845],
 [ 43060337917066, 12505219714880, 1012046350098],
 [ 44042799691375, 12765899584000, 1034164628272],
 [ 45001062940950, 13011196735808, 1052981571296],
 [ 45950826837754, 13279380818496, 1073830116063],
 [ 46905094973278, 13545048903104, 1094524279796],
 [ 47878639010788, 13804359927616, 1115702682066],
 [ 48846853321243, 14047296410048, 1135233704355],
 [ 49802458502478, 14313996702528, 1156012315681],
 [ 50773630761852, 14577302469440, 1177463371208],
 [ 51750458805180, 14835774389312, 1198617815583],
 [ 52739478013711, 15080810576128, 1219150608258],
 [ 53697957020474, 15345714552960, 1239822813201],
 [ 54654841695832, 15610873138816, 1260417321873],
 [ 55626978268439, 15871184033472, 1281344717506],
 [ 56608356014892, 16128617872000, 1302888854736],
 [ 57596587582005, 16385920936832, 1324757509346],
 [ 58972704627927, 16386082136704, 1324757509346],
 [ 60347539810597, 16386234461888, 1324757509346],
 ...])
```

- These three columns represent the three core counters we count: SSE_FLOPS, Data Cache Sys Fills, and User Cycles
- There is one row for each time sample
- How do we know which column corresponds to which counter?
- The get_schema method returns the schema for this data for a given initial key

```
In [8]: j.get_schema('amd64_core')
```

```
Out[8]: {'USER': (is_event=True), 'DCSF': (is_event=True, unit=B), 'SSE_FLOPS': (is_event=True)}
```

- The keys of this dictionary let you know what's available
- There are some additional things like the unit of measure for a thing, and whether it's an event or not
- To get the column index into the host stats array, you can do:

```
In [9]: index=j.get_schema('amd64_core')['SSE_FLOPS'].index
        print index
```

```
2
```

- So now we can do:

```
In [10]: v=j.hosts[h].stats['amd64_core']['0'][:,index]
        print v
```


```
[
  0      23836512957   57033734781   72806955413   94484478062
116993344068 140416884014 164308891339 184305633969 205218923263
226568975083 249877419073 270527680112 291467237077 313293140007
336508041477 357992006775 379178111515 401003272674 424697788352
447578226683 467208752172 488585464240 511800133446 533044554937
554174603954 576288959216 599796891201 620612696985 641902657767
665094588944 686540868604 708395167192 730599101169 752778153174
774810459287 796211334444 819180713197 840090212933 861444700425
882681660275 905701392186 928885469720 951572953447 970512446915
991327331845 1012046350098 1034164628272 1052981571296 1073830116063
1094524279796 1115702682066 1135233704355 1156012315681 1177463371208
1198617815583 1219150608258 1239822813201 1260417321873 1281344717506
1302888854736 1324757509346 1324757509346 1324757509346 1324757509346
1324757509346 1324757509346 1324757509346 1324757509346 1324757509346
1324757509346 1324757509346 1324757509346 1324757509346 1324757509346
1324757509346 1324757509346 1324757509346 1324757509346 1324757509346
1324757509346 1324757509349 1324757509349 1324757509349 1324757509349
1324757509349 1324757509349 1324757509349 1324757509349 1324757509349
1324757509349 1324757509349 1324757509350 1324757509350 1324757509350
1324757509350 1324757509350 1324757509350 1324757509350 1324757509350
1324757509350 1324757509350 1324757509350 1324757509350 1324757509350
1324757509350 1324757509350 1324757509350 1324757509350 1324757509350
1324757509350 1324757509350 1324757509350 1324757509350 1324757509350
1324757509350 1324757509350 1324757509350 1324757509350 1324757509352
1324757509352 1324757509352 1324757509352 1324757509352 1324757509352
1324757509356 1324757509356 1324757509356 1324757509356 1324757509356
1324758011074]
```

- Which is a 1-D numpy array of values for SSE_FLOPS for core 0 on host i101-111 on Ranger

- ```
In [11]: import numpy
 r = numpy.diff(v)/numpy.diff(j.times)
 print r
```

```
In [12]: %pylab inline
 from pylab import *
 t=j.times-j.times[0]
 ax=subplot(111)
 plot(t[0:-1]/3600.,r)
```

```
Out[12]: []
```



```
In [13]: help(j.hosts[h].get_stats)
```

Help on method get\_stats in module job\_stats:

```
get_stats(self, type_name, dev_name, key_name) method of job_stats.Host instance
Host.get_stats(type_name, dev_name, key_name)
Return the vector of stats for the given type, dev, and key.
```

```
In [14]: j.hosts[h].get_stats('amd64_core','0','SSE_FLOPS')
```

```
Out[14]: array([
 0, 23836512957, 57033734781, 72806955413,
 94484478062, 116993344068, 140416884014, 164308891339,
 184305633969, 205218923263, 226568975083, 249877419073,
 270527680112, 291467237077, 313293140007, 336508041477,
 357992006775, 379178111515, 401003272674, 424697788352,
 447578226683, 467208752172, 488585464240, 511800133446,
 533044554937, 554174603954, 576288959216, 599796891201,
 620612696985, 641902657767, 665094588944, 686540868604,
 708395167192, 730599101169, 752778153174, 774810459287,
 796211334444, 819180713197, 840090212933, 861444700425,
 882681660275, 905701392186, 928885469720, 951572953447,
 970512446915, 991327331845, 1012046350098, 1034164628272,
1052981571296, 1073830116063, 1094524279796, 1115702682066,
1135233704355, 1156012315681, 1177463371208, 1198617815583,
1219150608258, 1239822813201, 1260417321873, 1281344717506,
1302888854736, 1324757509346, 1324757509346, 1324757509346,
1324757509346, 1324757509346, 1324757509346, 1324757509346,
1324757509346, 1324757509346, 1324757509346, 1324757509346,
1324757509346, 1324757509346, 1324757509346, 1324757509346,
1324757509346, 1324757509349, 1324757509349, 1324757509349,
1324757509349, 1324757509349, 1324757509349, 1324757509349,
1324757509349, 1324757509349, 1324757509349, 1324757509349,
1324757509349, 1324757509350, 1324757509350, 1324757509350,
1324757509350, 1324757509350, 1324757509350, 1324757509350,
1324757509350, 1324757509350, 1324757509350, 1324757509350,
1324757509350, 1324757509350, 1324757509350, 1324757509350,
1324757509350, 1324757509350, 1324757509350, 1324757509350,
1324757509350, 1324757509350, 1324757509352, 1324757509352,
1324757509352, 1324757509352, 1324757509352, 1324757509352,
1324757509356, 1324757509356, 1324757509356, 1324757509356,
1324757509356, 1324758011074], dtype=uint64)
```

```
j.get_stats('amd64_core','0','SSE_FLOPS')
```

```
In [18]: j.acct['account']
```

```
Out[18]: 'TG-DMR090026'
```

```
In [19]: j.acct.keys()
```

```
Out[19]: ['ru_isrss',
 'ar_submission_time',
 'ru_inblock',
 'io',
 'owner',
 'slots',
 'id',
 'category',
 'queue',
 'ru_oublock',
 'group',
 'exit_status',
 'ru_stime',
 'ru_nsignals',
 'hostname',
 'arid',
 'priority',
 'failed',
 'ru_ixrss',
 'ru_nivcsw',
 'department',
 'pe_taskid',
 'task_number',
 'granted_pe',
 'mem',
 'start_time',
 'ru_msgsnd',
 'ru_wallclock',
 'ru_minflt',
 'submission_time',
 'maxvmem',
 'ru_nvcsw',
 'ru_nswap',
 'ru_majflt',
 'account',
 'iow',
 'name',
 'project',
 'ru_utime',
 'ru_ismrss',
 'end_time',
 'ru_idrss',
 'ru_maxrss',
 'ru_msgrcv',
 'cpu']
```

```
In [20]: j.acct['submission_time']
```

```
Out[20]: 1351525419
```

```
In []:
```